

Pricing the AI's Contribution, Not the AI: Evidence from the Audiobook Market

Yu Pan¹

¹Department of Business Administration, University of Zurich, Plattenstrasse 14, 8032 Zurich, Switzerland
Correspondence: yu.pan@business.uzh.ch

Abstract

Amazon's Virtual Voice program created the first mass market in which AI narrates human-written books under a label and each title, whether human- or AI-narrated, is rated separately on the dimension the narrator performs (Performance) and the dimension the author writes (Story). I exploit this attribute-split rating system in a census of Audible's 672,374 English-language audiobooks, 114,553 of them AI-narrated. The market's response is not a blanket rejection: the discount lands almost entirely on Performance (-0.65 stars with controls) while Story is statistically indistinguishable from zero, and the same split holds within author (-0.72 on Performance, -0.07 on Story). The summary Overall rating falls by less than the two component ratings imply; conditional on them, AI narration carries no further discount. In the small licensed-clone channel, where professional narrators lend their voices under contract, early evidence finds a far smaller discount, suggesting it attaches to the anonymous channel rather than to AI as such. AI narration reached 48% of new 2025 English-language releases, most of it from authors without an earlier human-narrated book; yet among 2024–25 releases only 0.1% of AI-narrated titles have reached 200 ratings, against 7.3% of human-narrated ones. Whatever mix of true quality and label response drives the Performance discount, it stops at the AI's attribute: the market prices the AI's contribution, not the AI.

Keywords generative AI, AI disclosure, voice cloning, online ratings, creative industries, algorithm aversion

Classification Social Sciences, Economic Sciences

Significance statement

Generative artificial intelligence is entering creative industries under disclosure labels, and Amazon's Virtual Voice program has created the first mass market in which labeled AI narration competes with human narration at scale. Audible rates every audiobook separately on the dimension the AI performs (narration) and the dimension the author writes (story), alongside an overall summary scale. Across a census of the store, the market's response concentrates almost entirely on the narration scale, leaves the writing scale essentially untouched, and adds no discount to the summary verdict beyond the narration component itself. Yet almost no AI-narrated title becomes widely rated. Labeled AI content is evaluated dimension by dimension: the market prices the AI's contribution, not the AI.

Introduction

In November 2023, Amazon began offering self-published authors a way to turn an e-book into an audiobook without hiring a narrator: a synthetic voice reads the text, and the title sells on Audible alongside human-narrated books (5; 21). By 2025, AI narration accounted for 47.7% of the platform's new English-language releases (Figure 1A), and every such title is credited to its synthetic narrator on the product page. Audiobooks are thus the first mass market in which AI performs human-authored work under a label, at scale, and where the market's own ratings score the narration separately from the writing. Whether such content arrives is no longer the open question. What remains open is where the market's response lands: on the product as a whole, or only on the AI's contribution to it.

Audible's rating system splits the evaluation along exactly the line the question requires. AI enters the product at a single step, the narration, while the writing remains the author's. Every title is rated separately, on the same five-star scale, for Performance, the narration, and for Story, the writing, both scored by the same audience for the same product. The AI supply also divides into two labeled channels: anonymous synthetic voices with no performer behind them, and licensed clones that carry a professional narrator's name. I take a census of the store and analyze its 672,374 English-language audiobooks, 114,553 of them AI-narrated.

Four patterns emerge. First, in counts, AI narration is additive entry: human-narrated output has not retreated, and most AI titles come from authors with no earlier human-narrated audiobook. Second, the evaluation is split, not blanket: with controls, AI-narrated books lose 0.65 stars on Performance, the

dimension the AI performs, while Story ratings move little or not at all, and the same split holds within author; nothing beyond the two components reaches the summary Overall rating. Third, the author's reputation does not blunt the discount, which tracks the channel instead, concentrating on anonymous voices, while licensed clones show, on early evidence, no comparable discount. Fourth, the flood is absorbed without breakout: AI-narrated books almost never become widely rated titles, while human-narrated books regularly do. In short, the market's evaluative response lands not on the product as a whole but on the AI's contribution to it.

One literature isolates the label. Since Dietvorst et al. (11) named algorithm aversion, experiments have held quality fixed, varied only what evaluators are told, and found resistance to AI provision that persists even at identical performance (19); in the closest antecedent to my setting, a blind preference for AI-generated content erodes once its origin is revealed (20). Whether that response governs a real market, at scale and with quality free to vary, is a question those designs cannot answer. The other literature reaches the market but not the structure of its response: LLM-written books, detected rather than disclosed, nearly tripled entry on Amazon while concentrating at the low-usage end of the market (22); labeled AI images expanded a stock-image marketplace while crowding out non-AI supply (14); on a large Chinese audio platform, AI entry shifted listening away from human-narrated titles (26); and generative AI cut demand for exposed freelance creative work (10). None of these settings rates the AI's contribution and the human's separately for the same product, and none observes a licensed channel that unbundles the technology from anonymity. The audiobook market has both, which is why it can show not only whether labeled AI content is discounted, but what the discount is a discount on.

My estimates are equilibrium measurements, in the descriptive tradition of large-scale studies of newly opened AI channels (16). Authors choose whether, and for which books, to adopt AI narration, so the estimates describe the market as it operates, not the effect of assigning a book to an AI voice. The Performance discount itself is a mixture: it bundles true quality differences with whatever response the visible label provokes, and my data do not separate the two. The paper's claims survive that ambiguity: wherever the discount comes from, what I establish is where it stops. Clone parity, likewise, is early evidence from a small, publisher-curated channel, and it bundles the narrator's own voice with loyal audiences and the curation itself. Ratings, finally, measure expressed evaluations, not demand or welfare. I interpret every result within these limits.

Data and Empirical Strategy

AI Narration on Audible

AI narration entered the audiobook market through Amazon's self-publishing arm. In November 2023, Kindle Direct Publishing opened Virtual Voice, an invitation-based program that converts an author's eligible e-book into an audiobook read by a generic synthetic voice at no production cost, an arrangement in place from launch through the census date. Authors set the list price and collect royalties, and the titles are sold on Audible,

Amazon's audiobook marketplace (5; 21). Disclosure is built in: each product page credits the narrator as "Virtual Voice" (hereafter, AI narration).

Two later programs professionalized the AI supply. In September 2024, Audible's production platform ACX began inviting professional narrators to license *voice replicas*, authorized AI copies of their own voices: rights holders select a replica for a title, the narrator reviews the production and is compensated per title, and the credit names the voice ("X's voice replica") (2; 7). In May 2025, Audible extended AI narration to selected publishers, offering more than one hundred synthetic voices across four languages (6). My analysis groups this supply into two channels that consumers can distinguish at the point of purchase: anonymous AI voices, meaning Virtual Voice and other synthetic voices that correspond to no performer, and licensed voice clones, which carry a professional narrator's name.

Audible's rating system is what makes this market a measurement instrument. Every title carries three rating scales: Overall, Performance, and Story. Performance covers the narration and Story the writing. A book's two dimension ratings average over the same audience, so the two are priced side by side for the same product, separating execution from the underlying work in a way a single overall rating cannot; the free-standing Overall scale returns in Section 3 as a check on the summary verdict. For an AI-narrated title, Performance then isolates the AI's contribution and Story the author's.

Data

I observe the complete Audible catalog as of July 2026, enumerated from the store's public sitemap and category listings. The *full sample* contains its 672,374 English-language audiobooks, 114,553 of them AI-narrated. For each title I observe product metadata and, for each rating scale, the cumulative average and the count of ratings; I do not observe individual ratings. The author is the book's first-listed writer, matched across narration regimes by a normalized name key and verified for authors appearing in both regimes (Supplementary Material, section S1).

Rating analyses use an *analysis sample* of books with at least five ratings on each dimension: 384,297 human-narrated and 14,933 AI-narrated books (Table 1). The table records the contrasts that motivate my empirical strategy: AI-narrated books are more recent, shorter, overwhelmingly self-published, mostly by authors without an earlier human-narrated book, and have far fewer ratings; the estimators in Section 2.3 are designed so that these observable differences do not pass for a rating discount. Market-level exhibits use the full sample; each exhibit states its sample.

Empirical Strategy

My estimates rest on two nested comparisons, each holding fixed a larger share of what differs between AI-narrated and human-narrated books.

First, the *within-book gap*. For book i in genre $g(i)$ and release year $t(i)$, I estimate

$$Gap_i = \beta AI_i + \mathbf{X}_i' \gamma + \mu_{g(i)} + \lambda_{t(i)} + \varepsilon_i, \quad (1)$$

Table 1. Summary Statistics, Analysis Sample

	Human-narrated				AI-narrated			
	Mean	Standard deviation	Minimum	Maximum	Mean	Standard deviation	Minimum	Maximum
Overall rating	4.42	0.41	1	5	4.28	0.51	1	5
Performance rating	4.50	0.41	1	5	3.88	0.63	1	5
Story rating	4.42	0.41	1	5	4.34	0.49	1	5
Ratings per book	335	2,274	5	306,227	20	34	5	2,022
Runtime (hours)	8.35	6.01	0	280	6.38	6.19	0	145
Number of narrators	1.28	1.06	1	16	1.00	0.01	1	2
Self-published	0.00	0.01	0	1	0.60	0.49	0	1
In a series	0.44	0.50	0	1	0.00	0.03	0	1
New author (indicator)	0.28	0.45	0	1	0.70	0.46	0	1
Books	384,297				14,933			
Authors	115,336				4,873			

Note. Analysis sample: English-language audiobooks on Audible as of July 2026 with at least five ratings on each dimension; ratings are cumulative averages over the book's audience. The author is the book's first-listed writer; the narrator is credited separately (a human performer or an AI voice). Self-published = publisher listed as Independently Published. New author = no human-narrated audiobook released before the given book. AI = Audible Virtual Voice (launched November 2023).

where Gap_i is the book's Performance rating minus its Story rating, AI_i indicates AI narration, and $\mathbf{X}_i$ collects log rating volume, runtime, a self-published indicator, and a series indicator. Both dimension ratings of a book aggregate essentially the same audience: within a book, the two rating counts are nearly identical, with a median difference of zero. Equation (1) is therefore equivalent to a book fixed-effects specification on the stacked dimension ratings. Factors that operate at the book level, including selection into the book's audience, that audience's overall generosity, vintage, price, genre appeal, and demand, all difference out. The controls and fixed effects in equation (1) therefore absorb only differences in the gap itself, such as genres or cohorts whose ratings tilt toward one dimension. I estimate the same specification with Performance and Story themselves as outcomes to show where the gap originates, and in the reported estimates I split the AI coefficient into its two channels.

Second, the *within-author discount*. For authors publishing in both narration regimes, I estimate

$$Y_i = \alpha_{a(i)} + \beta AI_i + \mathbf{X}_i' \boldsymbol{\gamma} + \lambda_{t(i)} + \varepsilon_i, \quad (2)$$

for each outcome $Y \in \{\text{Performance}, \text{Story}, \text{Gap}\}$, where the author fixed effect $\alpha_{a(i)}$ holds the author's identity fixed while the narration regime varies. Standard errors are clustered by author throughout.

These comparisons trace the structure of the market's response, not its composition: the Performance discount they recover is a mixture of genuine quality differences and reactions to the visible label, and my data do not separate the two. What the design does establish is where the response lands and where it does not.

All analyses were conducted in Python 3.14 with pandas 3.0.1, statsmodels 0.14.6, and pyfixest 0.50.1; figures were produced with matplotlib 3.10.8. The author used Claude Code (Anthropic) to assist with aspects of data analysis, programming, and manuscript editing. All outputs generated by these tools were carefully reviewed and validated by the author. The author remains solely responsible for the accuracy, integrity, and content of the work.

Results

The Flood: Entry Without Narrators

Figure 1A shows the inflow. AI narration launched in November 2023 and reached nearly half of new releases within two years: 4,029 AI-narrated titles in 2023 (10.1% of new English-language releases), 18,174 in 2024 (33.8%), and 33,869 in 2025 (47.7%). Human-narrated output did not retreat: roughly 36,000 titles per year across 2023–2025, in line with 2022. In counts, AI narration is additive entry, not displacement.

Figure 1B shows where the inflow concentrates. AI narration reaches all twelve of the largest categories, at depths ranging from 24.9% of new Science Fiction & Fantasy releases to 65.7% of Religion & Spirituality, with Erotica at 57.5%. Literature & Fiction, the largest category, sits at 41.0% and alone contributes 21.6% of AI-narrated output.

The Bill: Consumers Price the AI's Contribution

Figure 2A shows unadjusted means on the analysis sample. Performance falls from 4.50 on human-narrated books to 3.88 on AI-narrated books; Story falls from 4.42 to 4.34. Before any control, the two dimensions diverge: a drop of 0.62 against one of 0.08.

Table 2 reports the regression estimates. Columns (1)–(3) estimate equation (1) across all books in the analysis sample, splitting the AI coefficient into its two channels. Anonymous AI voices bear a Performance discount of 0.651 stars, a Story coefficient of -0.015 that is statistically indistinguishable from zero, and a within-book gap 0.636 lower. Licensed voice clones show no coefficient significant at the five percent level on any dimension; Section 3.3 returns to them.

Columns (4)–(6) estimate equation (2) on a sharper comparison: they keep only the 1,036 authors who publish both human-narrated and anonymous AI titles and add author fixed effects, so each author's AI-narrated books are compared with the same author's human-narrated books. Within these authors,




Figure 1 New Releases by Narration Regime, 2015–2025, and AI Share by Category, 2024–2025

(A) New English-language audiobook releases by release year, full sample (the complete Audible catalog as of July 2026); percentages are the AI share of each year's releases. Virtual Voice launched in November 2023; 2026 is omitted as a partial year. Titles are counted as live at the census date, so earlier vintages are lower bounds and recent AI shares are approximate. (B) AI share of 2024–2025 releases within the twelve largest categories by release volume. Categories are each title's primary store category. AI = Audible Virtual Voice.

Alt text: Two-panel chart. Panel A shows stacked bars of new audiobook releases per year from 2015 to 2025, in thousands; AI-narrated releases appear in 2023 and reach 47.7 percent of 2025 releases, while human-narrated counts stay near thirty-five thousand per year. Panel B shows horizontal bars of the AI share of 2024 to 2025 releases in the twelve largest categories, from 24.9 percent for Science Fiction and Fantasy to 65.7 percent for Religion and Spirituality.

an AI-narrated book is rated 0.724 lower on Performance and a statistically significant 0.067 lower on Story; the gap estimate of -0.657 is close to its pooled counterpart. Were the discount driven by weaker manuscripts routed to AI narration, Story, the dimension that measures the writing, would fall by more than Performance rather than by roughly a tenth as much. Manuscript writing quality thus explains little of the Performance discount. The diagnostic speaks to the writing; it does not rule out selection on unobserved characteristics that matter only for narration. Whichever of mild spillover or residual selection accounts for the remaining Story movement, roughly nine-tenths of the discount sits on the dimension the AI performs. The discount also remains large, and the Performance–Story gap essentially unchanged, when the inclusion threshold rises from five to fifty ratings on each dimension, so containment is not an artifact of the sample cutoff (Supplementary Material, section S3). The discount concentrates on the dimension the AI performs; the author's dimension moves little or not at all.

The Overall scale provides a third reading (Figure 3). On human-narrated books the summary verdict is well approximated by a weighted average of the two components, with Story carrying about 0.69 of the weight and Performance about 0.27: the market's summary verdict weighs the writing far more than the narration. Applied to the component estimates above, those weights imply an Overall discount of 0.19 stars; the estimated discount is 0.10. Conditional on both components, AI narration carries no further Overall discount: the residual term is slightly positive, in the regression and in a nonparametric comparison inside 0.1-star Performance $\times$ Story cells (Supplementary Material, section S5). Whatever the response to AI narration contains, the part that reaches the summary verdict arrives through the component the AI performs, not around it.

Who Bears the Discount

An AI-narrated book can carry two names: the author's on the cover and, for licensed clones, the narrator's on the credit. Table 2 locates the discount on the books that carry only the first: anonymous AI voices bear it; licensed voice clones largely do not. The two channels also draw on different suppliers. Of the 114,553 AI-narrated titles, anonymous voices account for 113,908 and licensed clones for 645. Among the anonymous titles, 64.9% are self-published and 21.7% are by authors with an earlier human-narrated audiobook; the clones are the mirror image, none self-published and 59.2% by such incumbent authors.

The author's reputation does not soften the discount. On human-narrated books the within-book gap is nearly flat in the author's prior human-narrated catalog (-0.003 per log point of prior ratings); on AI-narrated books it widens roughly ten times faster (-0.029 ; Supplementary Material, section S2): reputation lifts Story ratings on these books, but not the AI's Performance. Established authors face, if anything, a larger discount.

The clone channel is small and young, and I read it as early evidence. Of its 645 titles, 71 enter the analysis sample; 24.3% of the 2025 clone cohort has reached five ratings on both dimensions, and only 11.7% of the 428 titles released in the first half of 2026 have been rated at all. With that caveat, the evidence points the same way: in the analysis sample, clone books average 4.56 on Performance against 4.58 on Story; Table 2 finds no discount significant at the five percent level on any dimension; and against the same publishers' human-narrated books the within-book gap differs by -0.068 (0.082) (note to Table 2). For the thirteen narrators with rated books both in their own voice and through a clone (224 books), the within-narrator Performance difference is -0.098 : imprecise, but the confidence interval rules out a discount of the anonymous-channel size (Supplementary Material, section S4). The near-parity bundles a carried-over voice, loyal audiences, and publisher curation; my data do not separate them.

Absorption Without Breakout

Figure 2B tracks the 2024–2025 release cohorts across cumulative-ratings thresholds. The new supply is absorbed: 60.4% of AI-narrated releases attract at least one rating, against 94.1% of human-narrated releases. It does not break out: 6.7%

Table 2. The AI-Narration Discount: Pooled and Within-Author Estimates

	All books			Authors publishing in both regimes		
	(1) Performance	(2) Story	(3) Performance–Story	(4) Performance	(5) Story	(6) Performance–Story
AI narration, anonymous voice	-0.651*** (0.019)	-0.015 (0.015)	-0.636*** (0.013)	-0.724*** (0.033)	-0.067*** (0.023)	-0.657*** (0.027)
AI narration, licensed voice clone	-0.104* (0.057)	-0.013 (0.035)	-0.091* (0.049)			
Book controls	Yes	Yes	Yes	Yes	Yes	Yes
Genre fixed effects	Yes	Yes	Yes	No	No	No
Release-year fixed effects	Yes	Yes	Yes	Yes	Yes	Yes
Author fixed effects	No	No	No	Yes	Yes	Yes
Mean of dependent variable	4.476	4.418	0.058	4.274	4.427	-0.154
Observations	399,230	399,230	399,230	14,187	14,187	14,187
Authors	119,162	119,162	119,162	1,036	1,036	1,036
Adjusted R^2	0.235	0.193	0.160			
Within R^2				0.296	0.067	0.344

Note. Outcomes are Performance and Story ratings on a one-to-five star scale. Each column is one regression on the analysis sample of Table 1. Columns (4)–(6) add author fixed effects and keep the 1,036 authors with both human-narrated and anonymous AI titles (11 further authors, whose only AI titles are clones, enter columns (1)–(3) only); no licensed-clone title enters these columns. Within a book the two dimension ratings average over the same audience (the two rating counts are nearly identical within title), so columns (3) and (6) equal the difference of the two dimension columns and net out any book-level factor. Book controls: log rating volume, runtime, self-published, series. Standard errors, clustered by author, in parentheses. Restricting columns (4)–(6) to identity-verified authors gives -0.684 (0.040) in column (6); with publisher fixed effects on the 576 titles of the 52 publishers that carry clone titles, the clone–human within-book gap is -0.068 (0.082). AI = Audible Virtual Voice; a licensed voice clone is an authorized replica of a professional narrator, an anonymous voice any other AI narration. * $p < .10$; ** $p < .05$; *** $p < .01$.

of AI-narrated books reach twenty ratings against 40.1% of human-narrated books; 1.7% against 22.2% reach fifty; and 0.1% against 7.3% reach two hundred. Rating counts fold in the propensity to rate as well as listening, so these are reach measures rather than usage; no plausible difference in that propensity, however, spans the seventy-fold gap at the two-hundred-rating threshold. Nearly half of the market's new English-language releases are AI-narrated (Figure 1A); almost none becomes a widely rated book.

Discussion

My study examines the first two and a half years of labeled AI narration on Audible, using the complete catalog and ratings that price the narration and the writing of each book side by side. The central pattern is one of containment: the discount is large, but it stops at the dimension the AI performs.

Experiments that hold quality fixed and vary only the label find aversion to AI provision even at identical performance (19); yet the same literature shows that the response is not uniform: it can favor the algorithm (17), it concentrates in subjective and hedonic domains (8; 18), and, within a single product, it attaches to symbolic rather than functional components (15). At census scale, the market's response has precisely this structure: the discount, whatever its mixture of a genuine quality gap and a reaction to the visible label, lands on the dimension the AI performs and leaves the author's dimension nearly untouched. A perception experiment on human and AI audiobook narration points the same way at the individual level: it finds human and AI narration rated nearly alike, and its authors attribute an audiobook's value chiefly to the text, treating the voice as a functional medium rather than the locus of artistic value (4). A label response that devalued the whole product would have

to appear in Story ratings at comparable size; Story moves by at most a tenth as much. It would likewise reach the summary Overall rating beyond what the two components imply; it does not (Figure 3). A response driven by the label as such should likewise surface among licensed clones, which are disclosed just as plainly; on early evidence it largely does not. The response behaves less like a blanket verdict on the technology and more like one confined to the dimension the AI performs.

The discount is not borne evenly. It falls on the anonymous channel: overwhelmingly self-published titles by authors without an earlier human-narrated audiobook, the suppliers whose entry the technology enabled in the first place. Reputation does not offset it: an author's name is a brand built on writing, and nothing in it certifies narration (1); the discount widens rather than narrows with the author's prior catalog. The result is a flood that is absorbed but does not break out. Earlier digitization shocks let cheap entry turn unpromising creative goods into surprise hits and enlarged the market's stock of winners (25; 3). The AI-narrated flood has not opened that channel, and early industry evidence from music and games suggests the same shape across creative markets (9; 13).

The licensed-clone channel points to where this market may adjust. A professional narrator's name is a reputational asset, and licensing lets that asset travel with the technology rather than against it (24): clone titles, though few, show at most a fraction of the anonymous channel's discount, and the channel is growing quickly, from three titles in 2024 to 214 in 2025 and 428 in the first half of 2026. The economics are familiar: superstar markets arise where technology extends a performer's limited personal reach (23), and a licensed clone relaxes the one constraint recording left in place, the narrator's own hours. Whether the near-parity reflects the carried-over voice, loyal audiences, or the publishers who curate the channel, my data




Figure 2 Mean Ratings by Dimension and Share of Books Reaching Cumulative-Ratings Thresholds, by Narration Regime

(A) Mean Performance and Story ratings, analysis sample of Table 1 (384,297 human-narrated and 14,933 AI-narrated books). Within a book, both dimension ratings average over the same audience. The y-axis is truncated (3.5–4.85) for readability. (B) Share of books whose rating count reaches each threshold, 2024–2025 release cohorts (72,653 human-narrated and 52,043 AI-narrated books). Counts are cumulative as of July 2026. AI = Audible Virtual Voice.

Alt text: Two-panel chart. Panel A shows paired bars of mean Performance and Story ratings for human-narrated and AI-narrated books: Performance 4.50 versus 3.88, Story 4.42 versus 4.34. Panel B shows paired bars of the share of 2024 to 2025 releases reaching twenty, fifty, and two hundred cumulative ratings; at two hundred, 7.3 percent of human-narrated books against 0.1 percent of AI-narrated books.

cannot say; at its current size, the channel supplies early evidence, not a verdict.

The pattern has implications on three fronts. For platforms, the instrument matters as much as the label: Audible pairs disclosure with ratings that price the AI’s contribution separately, and that pairing is what makes the market’s response legible; as the transparency obligations of the EU AI Act take effect in August 2026 (12), attribute-level feedback is a natural complement to disclosure. For publishers and authors, the results on who bears the discount locate the lever: what travels with a title is not the author’s name but the narrator’s, and licensing an established voice appears, on early evidence, to restore much of what an author’s reputation cannot. For regulators, the finding tempers the algorithm-aversion evidence




Figure 3 The Summary Scale: Overall-Rating Estimates and Their Component Account

Point estimates with 95% confidence intervals for the Overall rating, analysis sample of Table 1, specification of Table 2. The implied discount applies the aggregation weights estimated on human-narrated books, Performance 0.269 (0.003) and Story 0.692 (0.003), $R^2 = 0.91$, to the Table 2 component estimates; it is a derived quantity and is drawn without an interval. The conditional AI terms condition on both component ratings: adding them to the pooled regression, comparing inside 0.1-star Performance $\times$ Story cells (human-narrated and anonymous-AI books; 99.7% of the anonymous arm matched), and adding author fixed effects on the sample of Table 2, columns (4)–(6). Licensed clones enter the regressions through their own indicator. Alternative designs: Supplementary Material, section S5. AI = Audible Virtual Voice.

Alt text: Coefficient plot with five rows and a dashed vertical line at zero. The estimated Overall discount is minus 0.10 stars; the discount implied by the two component ratings is minus 0.19. Three conditional AI terms, from the regression, matched-cell, and within-author designs, sit between plus 0.08 and plus 0.09 with narrow confidence intervals.

that shapes expectations about AI labeling: disclosure in this market produced no blanket rejection, but a precise, attribute-level response.

My limitations follow from the design. The census is an equilibrium snapshot: authors decide which books get an AI voice, and nothing in the data is randomized. The within-author comparison holds the author fixed, and the Story dimension bounds what manuscript selection can explain. My estimates therefore describe the market as it operates, not what would happen if a given book were switched to an AI voice. That within-author comparison rests on the incumbent authors who publish in both regimes, a minority in a market whose AI titles come mostly from new authors, for whom the pooled estimate and the Story diagnostic carry the weight. A book also enters the rating analysis only after five ratings on each dimension, a bar AI-narrated titles clear less often, so the analyzed books are more heavily rated than the catalog as a whole. The catalog is observed live in July 2026 and covers the English-language market, so delisted titles are invisible, earlier vintages are undercounted, non-English AI narration is out of scope, and each book’s ratings arrive as a lifetime average rather than a time path. And ratings are evaluations volunteered by consumers who chose the book: they show how a self-selected audience prices the two contributions, not demand, prices paid, or welfare.

The design anticipates its second wave. A rating-only rescan of the same universe, scheduled for August 2026, turns the census into a panel. It will show whether the Performance discount persists as cohorts age, give a first direct look at survival and delisting, and measure how absorption behaves as AI entry keeps climbing. A weekly crawl of new releases adds a forward entry series untouched by survivorship, and the rescan will also show whether clone parity survives the channel's growth. If it does, the pattern of this market's first years would sharpen into a division of labor, anonymous AI absorbed at the low end and licensed identity carrying value through, and the scarce input in a market suddenly full of voices would be not narration but a name.

Funding

The author declares no funding.

Competing interest

The author declares no competing interests.

Data availability

The data underlying this article were collected from publicly accessible Audible catalog pages and contain book-level aggregates only. The full replication package, the collection and analysis code together with the book-level datasets, will be deposited in the Open Science Framework (OSF) upon acceptance.

References

- David A. Aaker and Kevin Lane Keller. Consumer evaluations of brand extensions. *Journal of Marketing*, 54(1):27–41, 1990.
- ACX. Now in beta: Narrator voice replicas on ACX, 2024. Accessed July 11, 2026, <https://www.acx.com/mp/blog/now-in-beta-narrator-voice-replicas-on-acx>.
- Luis Aguiar and Joel Waldfogel. Quality predictability and the welfare benefits from new products: Evidence from the digitization of recorded music. *Journal of Political Economy*, 126(2):492–524, 2018.
- Annika Ahrens-Schwabe, Lukas Kosch, Günther Stocker, and Hajo G. Boomgaarden. Who cares about artificial intelligence? Human and artificial voices in audiobooks. *Computers in Human Behavior Reports*, 22:101068, 2026.
- Amazon Kindle Direct Publishing. Beta: Audiobooks with virtual voice, 2023. Accessed July 11, 2026, https://kdp.amazon.com/en_US/help/topic/GMPQGZAZJH6FF456.
- Audible. Audible expands catalog with AI narration and translation for publishers, 2025. Accessed July 11, 2026, <https://www.audible.com/about/newsroom/audible-expands-catalog-with-ai-narration-and-translation-for-publishers>.
- Bloomberg News. Amazon's Audible to create AI voice replicas of select audiobook narrators. Bloomberg (September 9), 2024. <https://www.bloomberg.com/news/articles/2024-09-09/amazon-s-audible-to-create-ai-voice-replicas-of-select-audiobook-narrators>.
- Noah Castelo, Maarten W. Bos, and Donald R. Lehmann. Task-dependent algorithm aversion. *Journal of Marketing Research*, 56(5):809–825, 2019.
- Deezer. AI-generated tracks represent 44% of new uploaded music, 2026. Accessed July 12, 2026, <https://newsroom-deezer.com/2026/04/ai-generated-tracks-represent-44-of-new-uploaded-music/>.
- Ozge Demirci, Jonas Hannane, and Xinrong Zhu. Who is AI replacing? The impact of generative AI on online freelancing platforms. *Management Science*, 71(10):8097–8108, 2025.
- Berkeley J. Dietvorst, Joseph P. Simmons, and Cade Massey. Algorithm aversion: People erroneously avoid algorithms after seeing them err. *Journal of Experimental Psychology: General*, 144(1):114–126, 2015.
- European Union. Regulation (EU) 2024/1689 (Artificial Intelligence Act), Article 50: Transparency obligations for providers and deployers of certain AI systems, 2024. Transparency obligations applicable from August 2, 2026. Accessed July 12, 2026, <https://artificialintelligenceact.eu/article/50/>.
- Game Oracle. AI disclosure and the reception of new releases on Steam, 2026. Analysis by Ross Burton of 9,879 titles released January–October 2025; as reported in PC Gamer, June 2026, <https://www.pcgamer.com/software/ai/data-analyst-finds-ai-stigma-on-steam-can-reduce-the-number-of-reviews-a-game-gets-by-around-53-percent-and-the-reviews-it-does-get-are-more-negative/>.
- Samuel G. Goldberg and H. Tai Lam. Generative AI & creative goods: Market expansion, crowd-out, and copyright, 2026. Conditionally accepted, *Quarterly Journal of Economics*.
- Armin Granulo, Christoph Fuchs, and Stefano Puntoni. Preference for human (vs. robotic) labor is stronger in symbolic consumption contexts. *Journal of Consumer Psychology*, 31(1):72–80, 2021.
- Maximilian Kaiser and Christian Schulze. Frontiers: ChatGPT referrals to e-commerce websites: How do LLMs compare against traditional channels? *Marketing Science*, 45(4):699–715, 2026.
- Jennifer M. Logg, Julia A. Minson, and Don A. Moore. Algorithm appreciation: People prefer algorithmic to human judgment. *Organizational Behavior and Human Decision Processes*, 151:90–103, 2019.
- Chiara Longoni and Luca Cian. Artificial intelligence in utilitarian vs. hedonic contexts: The “Word-of-Machine” effect. *Journal of Marketing*, 86(1):91–108, 2022.
- Chiara Longoni, Andrea Bonezzi, and Carey K. Morewedge. Resistance to medical artificial intelligence. *Journal of Consumer Research*, 46(4):629–650, 2019.
- Petr Parshakov, Iuliia Naidenova, Sofia Paklina, Nikita Matkin, and Cornel Nesseler. Users favor LLM-generated content—until they know it's AI, 2025. arXiv preprint arXiv:2503.16458.
- Publishers Weekly. Kindle Direct Publishing will beta test virtual voice-narrated audiobooks. Publishers Weekly (November 1), 2023. <https://www.publishersweekly.com/pw/by-topic/industry-news/audio-books/article>

-
- le/93625-kindle-direct-publishing-will-beta-test-virtual-voice-narrated-audiobooks.html.
22. Imke Reimers and Joel Waldfogel. AI and the quantity and quality of creative products: Have LLMs boosted creation of valuable books? NBER Working Paper 34777, National Bureau of Economic Research, 2026.
 23. Sherwin Rosen. The economics of superstars. *American Economic Review*, 71(5):845–858, 1981.
 24. Steven Tadelis. What's in a name? Reputation as a tradeable asset. *American Economic Review*, 89(3):548–563, 1999.
 25. Joel Waldfogel. How digitization has created a golden age of music, movies, books, and television. *Journal of Economic Perspectives*, 31(3):195–214, 2017.
 26. Ning Wan, Yingjie Zhang, and Ting Li. When AI speaks, do humans lose listeners? Evidence from the audiobook industry, 2024. SSRN working paper 5062901.

Supplementary material for “Pricing the AI’s Contribution, Not the AI: Evidence from the Audiobook Market”

Yu Pan¹

¹Department of Business Administration, University of Zurich, Plattenstrasse 14, 8032 Zurich, Switzerland
Correspondence: yu.pan@business.uzh.ch

Abstract

This supplement collects the supporting analyses referenced in the main text: author matching and identity verification (section S1), reputation gradients by narration regime (section S2), the robustness of the containment result to the rating threshold (section S3), the classification of the licensed-clone channel together with within-narrator estimates (section S4), and the robustness of the summary-scale estimate (section S5). References of the form “Table 2,” “Figure 3,” or “equation (2)” refer to the main text; tables numbered S1–S5 are in this supplement.

S1. Author Matching and Identity Verification

The within-author estimator of equation (2) in the main text, and the reputation gradients of section S2, both link an author’s AI-narrated and human-narrated books. I match books to authors by a normalized name key: the first-listed author’s name, lowercased and stripped of punctuation and spacing. Because two distinct authors can share a name, the key can create false incumbents, and Table S1 verifies identity for the 4,393 keys that appear in both narration regimes.

Panel A classifies these keys by direct cross-regime evidence. Confirmed keys share, across the two regimes, an Audible author-page identifier, a series, or a normalized title (typically a dual edition of the same book): links that a shared name alone does not create. Likely keys share a distinctive publisher. Suspect keys are well covered by author-page identifiers on both sides with no overlap, or accumulate three or more distinct author pages without any bridge. The large Ambiguous class records the absence of identifiers rather than evidence of a mismatch: AI-narrated self-published titles usually carry no author-page identifier that could confirm or refute the match.

A false match staples one person’s AI-narrated books to another person’s human-narrated catalog, which dilutes incumbent-based contrasts toward the new-author baseline; gating on verified identity is therefore the conservative direction of the check. Panel B re-estimates the reputation gradient on AI-narrated books from Table S2, column (6), under the gates. The gradient is essentially unchanged in every gated sample, including the strictest, which keeps only new authors and Confirmed incumbents and removes 85% of incumbent books. The paper’s central within-author estimate behaves the same way: restricting the estimator to identity-verified authors gives a Performance–Story gap of -0.684 (0.040), against -0.657 in the full within-author sample (Table 2, column 6, main text),

on 6,138 books by 292 of the 1,036 authors in that sample. False name collisions therefore do not drive the within-author discount.

S2. Reputation Gradients by Narration Regime

Section 3.3 of the main text reports that an author’s reputation does not soften the AI-narration discount. Table S2 gives the underlying regressions. For each narration regime separately, I regress the book-level Performance rating, Story rating, and their difference on the author’s prior human-narrated reputation (the log of one plus the cumulative rating count on the author’s earlier human-narrated audiobooks at the book’s release), together with genre and release-year fixed effects and the book controls of the main text.

On human-narrated books the within-book Performance–Story gap is nearly flat in prior reputation (-0.003 per log point, column 3); on AI-narrated books it widens about ten times faster (-0.029 , column 6). The asymmetry has a clear source: reputation lifts Story ratings on an author’s AI-narrated books (column 5, $+0.018$) but not their Performance (column 4, -0.011 , not significant). An established author’s brand, built on writing, raises the writing score but does not transfer to the AI’s narration. In the identity-verified subsamples of section S1 the gap gradient is stable; the split between the two dimension slopes is estimated less precisely there, and the main text rests no claim on that split.

Table S1. Author-Identity Verification: Classification and Gated Re-Estimates*Panel A: cross-regime author keys by identity class*

	Author keys	Incumbent AI books
Confirmed: shared author page, series, or title	339	656
Likely: shared distinctive publisher	718	945
Suspect: conflicting author pages, no bridge	78	48
Ambiguous: identifiers missing on one side	3,258	2,839
All cross-regime keys	4,393	4,488

Panel B: reputation gradient on AI-narrated books, under identity gates

Sample	Gradient	(SE)	AI books	Incumbent books	Incumbent authors
Full analysis sample (baseline)	-0.029***	(0.006)	14,933	4,488	1,145
New authors + Confirmed incumbents	-0.031**	(0.014)	11,101	656	92
New authors + Confirmed or Likely	-0.027***	(0.009)	12,046	1,601	324
All classes except Suspect	-0.029***	(0.006)	14,885	4,440	1,132

Note. The author key is the normalized name of the book's first-listed writer. Panel A classifies the 4,393 keys with books in both narration regimes by direct cross-regime evidence. Confirmed: the two regimes share an author-page identifier, a series, or a normalized title. Likely: the two regimes share a distinctive publisher (at most 1,000 census titles, excluding Independently Published). Suspect: author pages are well covered on both sides with no overlap, or the key accumulates three or more distinct author pages without any bridge. Ambiguous: the remaining keys, on which AI-narrated titles carry no author-page identifier. Incumbent AI books = the class's AI-narrated books in the analysis sample (at least five ratings on each dimension) whose author has an earlier human-narrated audiobook. Panel B re-estimates the reputation gradient of Table S2, column (6), specification unchanged, on the AI-narrated books of the stated sample. New authors are retained throughout and the gates act on incumbent books only. Standard errors, clustered by author, in parentheses. AI = Audible Virtual Voice. * $p < .10$; ** $p < .05$; *** $p < .01$.

Table S2. Reputation Gradients by Narration Regime and Rating Dimension

	Human-narrated			AI-narrated		
	(1) Performance	(2) Story	(3) Performance–Story	(4) Performance	(5) Story	(6) Performance–Story
Prior human catalog, $\log(1 + \text{ratings})$	0.012*** (0.002)	0.015*** (0.002)	-0.003*** (0.001)	-0.011 (0.010)	0.018*** (0.007)	-0.029*** (0.005)
Genre fixed effects	Yes	Yes	Yes	Yes	Yes	Yes
Release-year fixed effects	Yes	Yes	Yes	Yes	Yes	Yes
Book controls	Yes	Yes	Yes	Yes	Yes	Yes
Mean of dependent variable	4.50	4.42	0.08	3.88	4.34	-0.46
Observations	384,297	384,297	384,297	14,933	14,933	14,933
Adjusted R^2	0.194	0.210	0.049	0.091	0.062	0.124

Note. Each column is one regression on the analysis sample (English-language audiobooks with at least five ratings on each dimension; book-level cumulative averages over the book's audience). The reputation measure is the log of one plus the cumulative rating count on the author's earlier human-narrated audiobooks at the book's release; every specification also includes an any-prior-catalog indicator, log rating volume, runtime, a self-published indicator, and a series indicator. Standard errors, clustered by author, in parentheses. Author identity is a normalized name key, verified in section S1. Columns (3) and (6) equal (1)–(2) and (4)–(5) by construction: within a book the same audience scores both dimensions, so audience generosity and selection difference out. AI = Audible Virtual Voice. * $p < .10$; ** $p < .05$; *** $p < .01$.

S3. Robustness of the Containment Result to the Rating Threshold

The analysis sample keeps books with at least five ratings on each dimension. Because AI-narrated titles clear that bar less often than human-narrated ones, one might worry that the containment pattern (a discount on Performance but not on Story) is an artifact of which books survive the cutoff. Table S3 re-estimates the pooled specification of Table 2 (columns 1–3, main text) at minimum-rating thresholds of five, ten, twenty, and fifty on each dimension.

Raising the threshold prunes AI-narrated books hardest, yet the anonymous-voice Performance discount stays large and significant at every cutoff, and the Performance–Story gap

barely moves (between -0.64 and -0.66). The Story coefficient is never negative; it drifts from zero at the baseline to mildly positive under the strictest cutoffs, consistent with positive selection of heavily rated AI titles on the writing dimension. Containment is thus not an artifact of the five-rating cutoff.

S4. The Licensed-Clone Channel: Classification and Within-Narrator Estimates

AI narration is classified into the two channels a consumer can tell apart at the point of purchase, read from the narrator credit itself. *Anonymous voices* are Virtual Voice and other synthetic voices credited to no performer. *Licensed voice clones* carry a professional narrator's name (for example, "X's voice replica"):

Table S3. Pooled Estimates at Alternative Minimum-Rating Thresholds

Minimum ratings per dimension	Performance		Story		Performance–Story		Books	AI
	Coef.	(SE)	Coef.	(SE)	Coef.	(SE)		
≥ 5 (<i>baseline</i>)	−0.651***	(0.019)	−0.015	(0.015)	−0.636***	(0.013)	399,230	14,862
≥ 10	−0.623***	(0.025)	0.028	(0.019)	−0.651***	(0.014)	321,001	8,409
≥ 20	−0.587***	(0.032)	0.074***	(0.022)	−0.660***	(0.019)	247,132	3,954
≥ 50	−0.541***	(0.059)	0.123***	(0.033)	−0.664***	(0.033)	160,954	964

Note. Outcomes are on a one-to-five star scale. Each row re-estimates the pooled specification of Table 2, columns (1)–(3), unchanged, on books with at least the stated number of ratings on each dimension. Cells report the anonymous-voice coefficient. Standard errors, clustered by author, in parentheses. The ≥ 5 row is the analysis sample used throughout the paper; it retains about 13% of AI-narrated but 69% of human-narrated books, and stricter cutoffs prune AI-narrated books faster still. The two rightmost columns give the total book count and the anonymous-voice AI book count (licensed clones enter the regressions through their own indicator and are excluded from that count). AI = Audible Virtual Voice. * $p < .10$; ** $p < .05$; *** $p < .01$.

Table S4. Within-Narrator Estimates: Licensed Clones Against the Same Narrator's Human-Narrated Titles

	Performance	Story	Performance–Story
Licensed voice clone	−0.098 (0.246)	0.154 (0.156)	−0.252 (0.206)
Narrator fixed effects	Yes	Yes	Yes
Release-year fixed effects	Yes	Yes	Yes
Book controls	Yes	Yes	Yes
Books	224	224	224
Narrators	13	13	13
Clone books	35	35	35

Note. Outcomes are on a one-to-five star scale. The sample contains every analysis-sample title (at least five ratings on each dimension) of the 13 professional narrators who record both in their own voice and through an authorized voice replica. Each column is one regression of the outcome on a clone indicator with narrator and release-year fixed effects; book controls are log rating volume, runtime, and a series indicator. Standard errors, clustered by narrator, in parentheses. AI = Audible Virtual Voice. * $p < .10$; ** $p < .05$; *** $p < .01$.

645 titles across 51 distinct voices, each carrying the replica wording in the credit, so this channel too is disclosed at the point of sale. The full sample partitions exactly into 113,908 anonymous and 645 clone titles; author-keyed analyses drop one AI title with no author key, so channel shares in the main text are computed on the 113,907 author-keyed anonymous titles.

Of the 51 cloned voices, 47 also narrate human-narrated audiobooks in the census (823 titles); 13 voices have titles passing the analysis-sample threshold on both sides, 224 books in all, 35 of them clones. Table S4 compares clone titles against the same narrator's human-narrated titles. The point estimates are imprecise on every dimension, and the Performance interval, though wide ($[-0.58, 0.38]$), excludes a discount of the anonymous-channel size (-0.651 pooled, -0.724 within author). With thirteen narrators the design has power only against large discounts; the table is reported for transparency and reads as consistent with the clone rows of Table 2 in the main text, not as an independent verdict.

Overall discount, the residual term being slightly positive. Table S5 re-estimates that term under alternative designs aimed at three concerns: the functional form and support of the aggregation (finer matching cells, a trim below the scale ceiling, deep common support, a quadratic aggregation function, and the common-support box of the AI arm), the composition wedge between Overall and dimension ratings (the subsample in which dimension ratings accompany at least 95% of Overall ratings), and unobserved author heterogeneity (author fixed effects on the within-author sample of Table 2). The estimate is positive and of similar size in every design, so the term does not reflect the aggregation's functional form, the bounded scale, or the wedge between the scales' audiences.

S5. Robustness of the Summary-Scale Estimate

Figure 3 of the main text reports that, conditional on the Performance and Story ratings, AI narration carries no further

Table S5. Robustness of the Summary-Scale Estimate

Specification	Coef.	(SE)	Observations
Baseline conditional regression (Figure 3)	0.083***	(0.004)	399,230
Cell fixed effects, 0.1-star cells (Figure 3)	0.091***	(0.002)	390,037
Cell fixed effects, 0.05-star cells	0.090***	(0.002)	356,135
Cell fixed effects, Overall < 4.9	0.091***	(0.002)	361,142
Cell fixed effects, Performance and Story < 4.6	0.097***	(0.003)	174,661
Common support of the AI arm	0.086***	(0.003)	289,967
Quadratic aggregation	0.091***	(0.004)	399,230
Dimension ratings accompany $\geq 95\%$ of Overall ratings	0.073***	(0.004)	104,178
Author fixed effects (Table 2 columns (4)–(6) sample)	0.089***	(0.010)	14,187

Note. Each row reports the anonymous-AI coefficient from a re-estimate of the summary-scale conditional specification (main text, Figure 3): the outcome is the Overall rating, and Performance and Story enter either as regressors (rows without cell fixed effects, with the Table 2 book controls and genre and release-year fixed effects) or through Performance $\times$ Story cell fixed effects (matching estimator rows, which drop licensed clones and keep cells containing both narration regimes). The common-support row keeps books inside the 5th–95th percentile box of the anonymous-AI arm's Performance and Story distributions. Standard errors, clustered by author, in parentheses. AI = Audible Virtual Voice. * $p < .10$; ** $p < .05$; *** $p < .01$.